

MODERN OPERATING SYSTEMS

Third Edition

ANDREW S. TANENBAUM

Chapter 13

Operating System Design

Operating System Goals

- Define abstractions.
- Provide primitive operations.
- Ensure isolation.
- Manage the hardware.

Difficulties Designing Operating Systems (1)

- Operating systems have become extremely large programs
- Must deal with concurrency
- Have to deal with potentially hostile users
- Many users want to share some of their information and resources with selected other users

Difficulties Designing Operating Systems (2)

- Operating systems live for a very long time
- Designers need to provide for considerable generality
- Systems are generally designed to be portable
- Operating systems need to be backward compatible with some previous operating system

Guiding Principles for Interface Design

- **Simplicity**

Perfection is reached not when there is no longer anything to add, but when there is no longer anything to take away. (Antoine de St. Exupery)

- **Completeness**

Everything should be as simple as possible, but no simpler. (Albert Einstein)

- **Efficiency**

If a feature or system call cannot be implemented efficiently, it is probably not worth having (Tanenbaum)

Execution Paradigms

```
main()  
{  
    int ... ;  
  
    init( );  
    do_something( );  
    read(...);  
    do_something_else( );  
    write(...);  
    keep_going( );  
    exit(0);  
}
```

(a)

```
main()  
{  
    mess_t msg;  
  
    init( );  
    while (get_message(&msg)) {  
        switch (msg.type) {  
            case 1: ... ;  
            case 2: ... ;  
            case 3: ... ;  
        }  
    }  
}
```

(b)

Figure 13-1. (a) Algorithmic code. (b) Event-driven code.

System Structure

- Layered Systems
- Exokernels
- Microkernel-Based Client-Server Systems
- Extensible Systems
- Kernel Threads

Layered Systems

Layer

7	System call handler					
6	File system 1		...		File system m	
5	Virtual memory					
4	Driver 1	Driver 2	...			Driver n
3	Threads, thread scheduling, thread synchronization					
2	Interrupt handling, context switching, MMU					
1	Hide the low-level hardware					

Figure 13-2. One possible design for a modern layered operating system.

Microkernel-Based Client-Server Systems

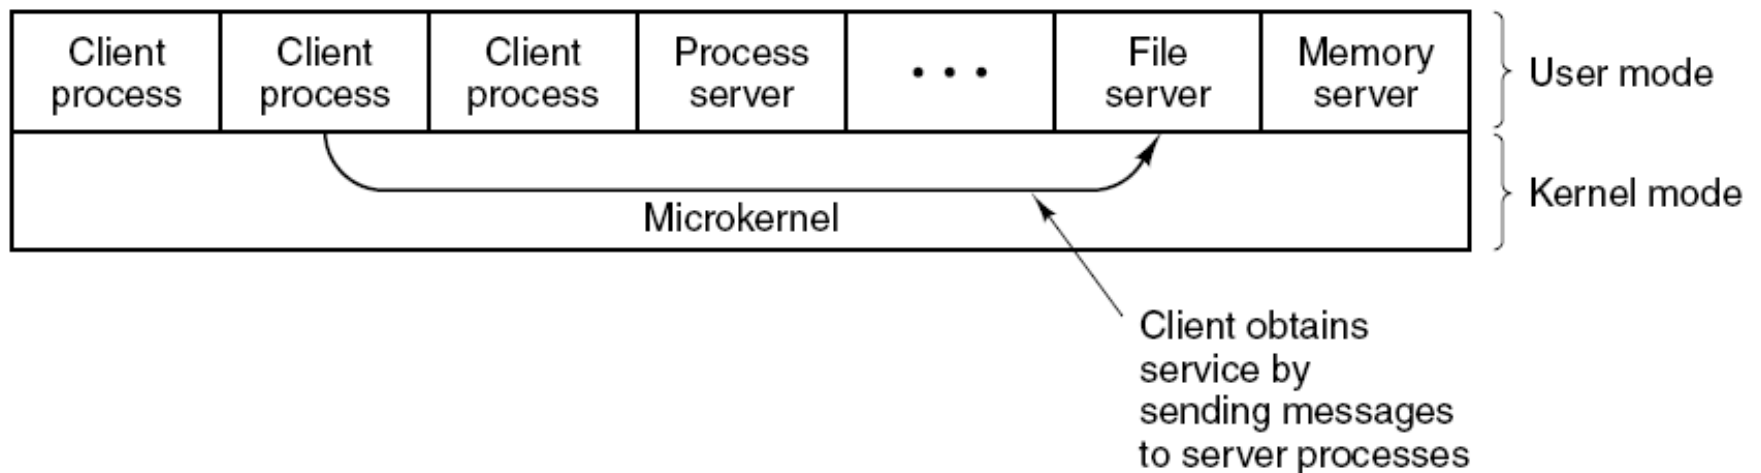


Figure 13-3. Client-server computing based on a microkernel.

Naming

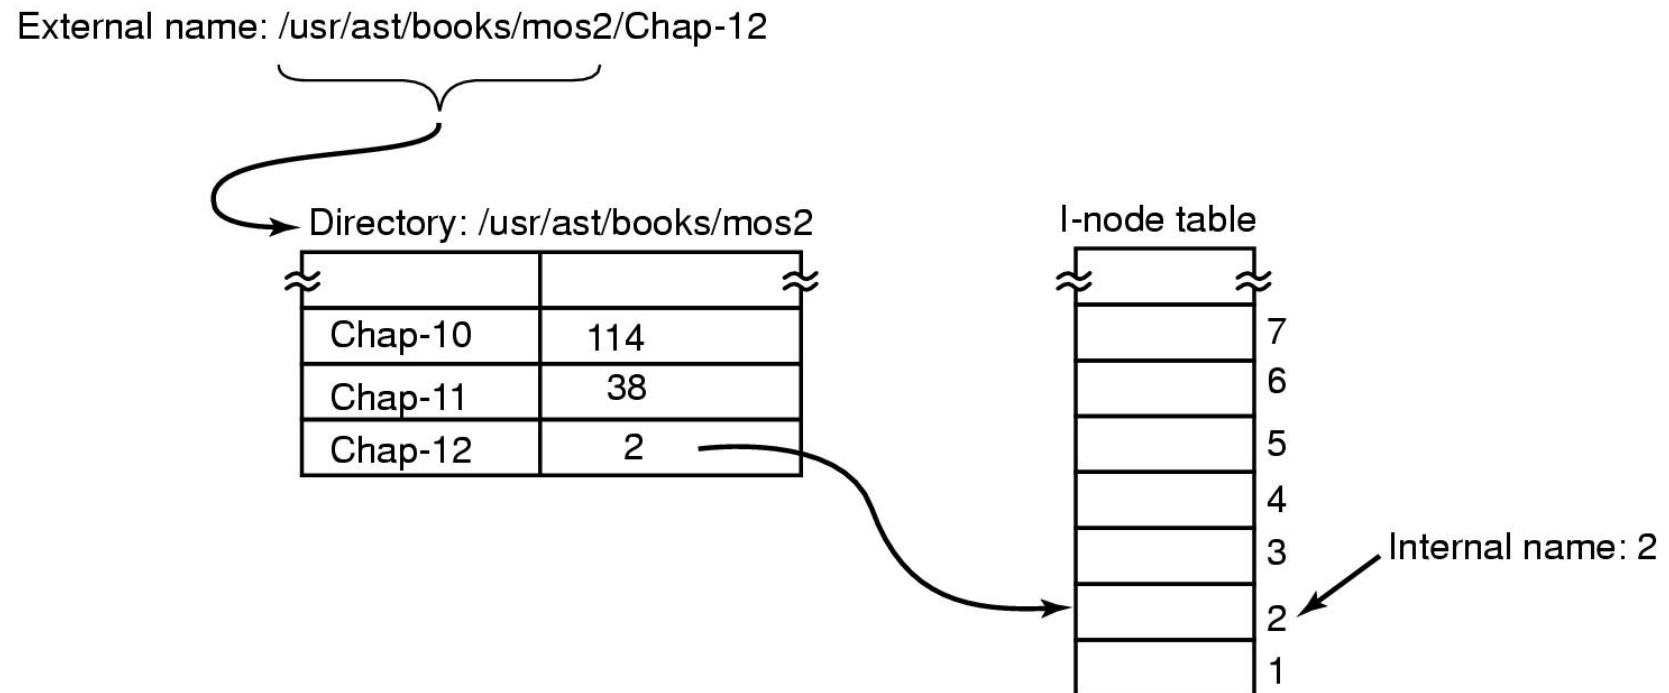


Figure 13-4. Directories are used to map external names onto internal names.

Binding Time

- Early binding
 - Simple
 - but less flexible
- Late binding
 - More complicated
 - but more flexible

Static versus Dynamic Structures

```
found = 0;
for (p = &proc_table[0]; p < &proc_table[PROC_TABLE_SIZE]; p++) {
    if (p->proc_pid == pid) {
        found = 1;
        break;
    }
}
```

Figure 13-5. Code for searching the process table for a given PID.

Useful Techniques

- Hiding the Hardware
- Indirection
- Reusability
- Reentrancy
- Brute Force
- Check for Errors First

Hiding the Hardware (1)

```
found = 0;
for (p = &proc_table[0]; p < &proc_table[PROC_TABLE_SIZE]; p++) {
    if (p->proc_pid == pid) {
        found = 1;
        break;
    }
}
```

Figure 13-6. (a) CPU-dependent conditional compilation. (

Hiding the Hardware (2)

```
#include "config.h"
```

```
init( )
```

```
#if (CPU = PENTIUM)
```

```
...
```

```
#endif
```

```
#if (CPU = SPARC)
```

```
...
```

```
#endif
```

```
#include "config.h"
```

```
#if (WORD_LENGTH == 32)
```

```
typedef int register
```

```
#endif
```

```
#if (WORD_LENGTH == 64)
```

```
typedef long register
```

```
#endif
```

Figure 13-6. (b) Word-length dependent conditional compilation.

Performance

- Why Are Operating Systems Slow?
- What Should Be Optimized?
- Space-Time Trade-offs
- Caching
- Hints
- Exploiting Locality
- Optimize the Common Case

Space-Time Trade-offs (1)

```
#define BYTE_SIZE 8                                /* A byte contains 8 bits */
int bit_count(int byte)
{
    int i, count = 0;
    for (i = 0; i < BYTE_SIZE; i++)
        if ((byte >> i) & 1) count++;
    return(count);
}
```

(a)

```
/*Macro to add up the bits in a byte and return the sum. */
#define bit_count(b) ((b&1) + ((b>>1)&1) + ((b>>2)&1) + ((b>>3)&1) + \
    ((b>>4)&1) + ((b>>5)&1) + ((b>>6)&1) + ((b>>7)&1))
```

(b)

```
/*Macro to look up the bit count in a table. */
char bits[256] = {0, 1, 1, 2, 1, 2, 2, 3, 1, 2, 2, 3, 2, 3, 3, 4, 1, 2, 2, 3, 2, 3, 3, ...};
#define bit_count(b) (int) bits[b]
```

(c)

Figure 13-7. (a) A procedure for counting bits in a byte.

(b) A macro to count the bits. (c) Table look up

Space-Time Trade-offs (2)

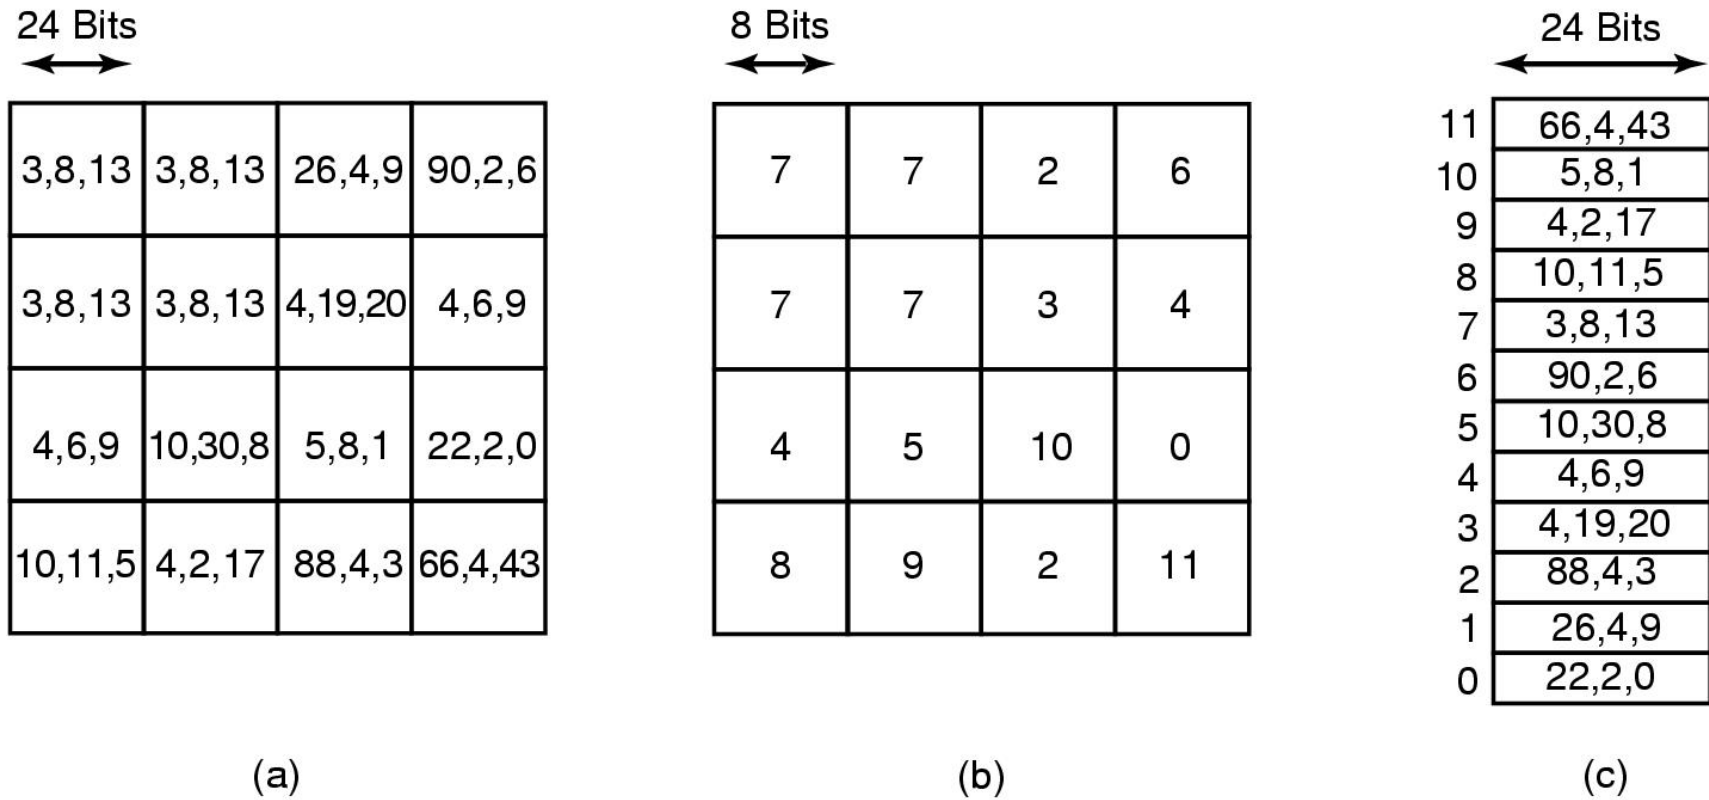


Figure 13-8. (a) Part of an uncompressed image with 24 bits per pixel. (b) With a palette

Caching (1)

To look up */usr/ast/mbox* (Fig. 4-35) requires the following disk accesses:

- Read the i-node for the root directory (i-node 1).
- Read the root directory (block 1).
- Read the i-node for */usr* (i-node 6).
- Read the */usr* directory (block 132).
- Read the i-node for */usr/ast* (i-node 26).
- Read *the /usr/ast* directory (block 406).

Caching (2)

Path	I-node number
/usr	6
/usr/ast	26
/usr/ast/mbox	60
/usr/ast/books	92
/usr/bal	45
/usr/bal/paper.ps	85

Figure 13-9. Part of the i-node cache for Fig. 4-35.

Project Management

The Mythical Man Month

- Large project design
 - 1/3 Planning
 - 1/6 Coding
 - 1/4 Module testing
 - 1/4 System testing
- People and time not interchangeable
 - Work cannot be fully parallelized
 - Work must be partitioned into large numbers of modules
 - Debugging is highly sequential

Team Structure

Title	Duties
Chief programmer	Performs the architectural design and writes the code
Copilot	Helps the chief programmer and serves as a sounding board
Administrator	Manages the people, budget, space, equipment, reporting, etc.
Editor	Edits the documentation, which must be written by the chief programmer
Secretaries	The administrator and editor each need a secretary
Program clerk	Maintains the code and documentation archives
Toolsmith	Provides any tools the chief programmer needs
Tester	Tests the chief programmer's code
Language lawyer	Part timer who can advise the chief programmer on the language

Figure 13-10. Mills' proposal for populating a 10-person chief programmer team.

The Role of Experience

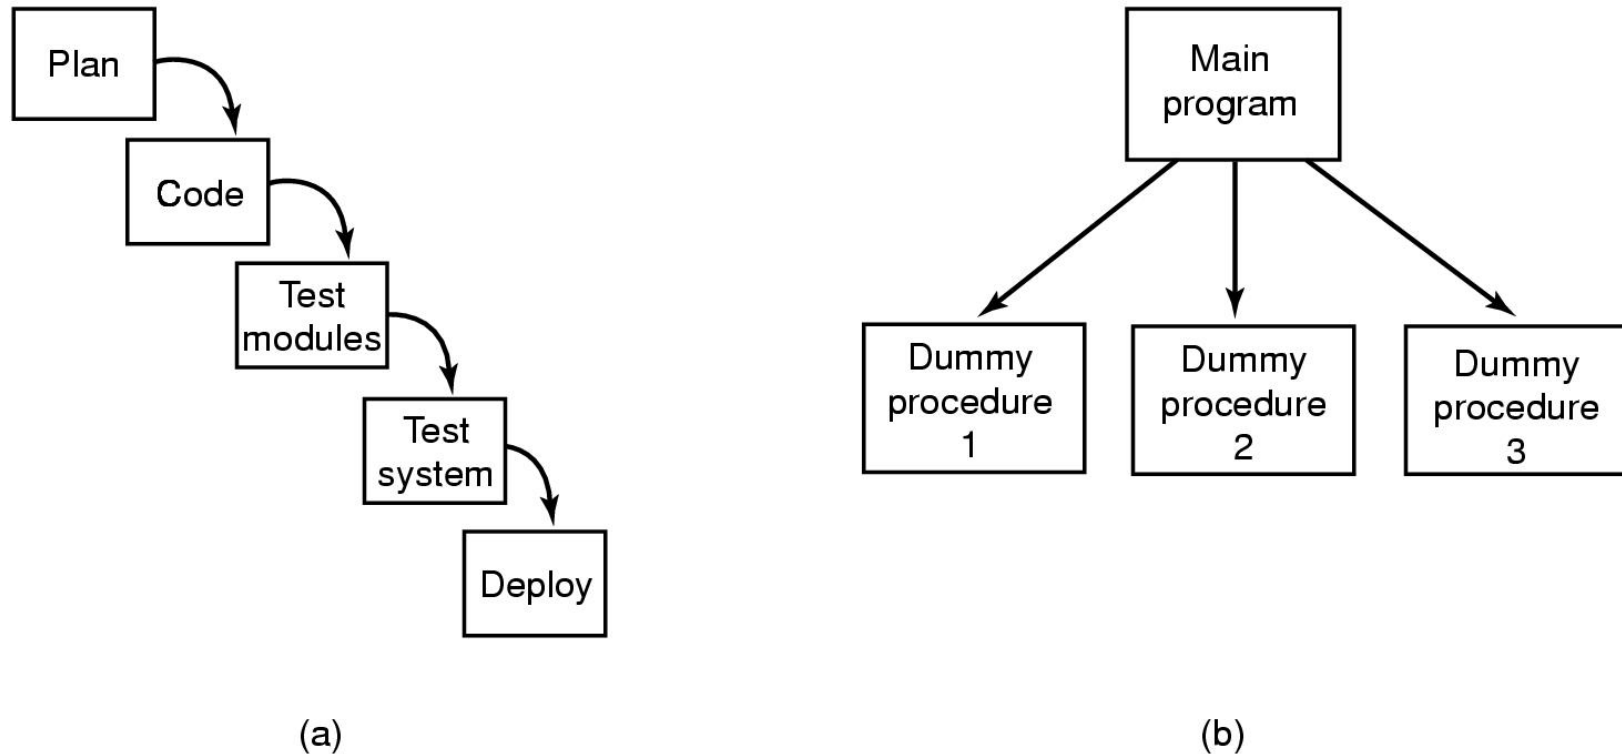


Figure 13-11. (a) Traditional software design progresses in stages.
(b) Alternative design produces a working system (that does nothing) starting on day 1.

Trends in Operating Systems

- Virtualization
- Multicore Chips
- Large Address Space Operating Systems
- Networking
- Parallel and Distributed Systems
- Multimedia
- Battery-Powered Computers
- Embedded Systems
- Sensor Nodes

Virtualization




Figure 13-12. A hypervisor running four virtual machines.